

The Process Control Block (PCB) is a data structure in the operating system kernel containing the information needed to manage the scheduling of a particular process. Here is a detailed breakdown of what a PCB typically contains:

Process Control Block (PCB)

A Process Control Block (PCB) is a data structure used by the operating system to manage information about a process. Here's a breakdown of the key components:

- **Process State:** The current state of the process (e.g., running, waiting, etc.).
- **Process ID (PID):** Unique identifier for the process.
- **Program Counter:** The address of the next instruction to be executed.
- **CPU Registers:** The contents of all process-centric registers, which are essential for the process's execution.
- **Memory Management Information:** Information about the process's memory allocation (e.g., page tables, segment tables), which help the operating system manage the process's memory usage.
- **Accounting Information:** Information about CPU usage, time limits, and other resource usage metrics.
- **I/O Status Information:** Information about I/O devices allocated to the process and the list of open files, enabling the operating system to manage the process's I/O operations effectively.

Simplified PCB Representation in C

Below is a simplified representation of a PCB in C:

```
struct PCB {  
    int process_id;           // Process ID  
    int process_state;        // Current state of the process
```

```
    unsigned int program_counter; // Address of the next instruction
    int cpu_registers[8];        // CPU registers (example with 8 registers)
    void *memory_base;          // Base address of the process's memory
    void *memory_limit;         // Limit address of the process's memory
    int priority;                // Priority of the process
    int accounting_info;         // Accounting information
    int io_status_info;          // I/O status information
};
```

This structure is used by the operating system to keep track of all the processes and manage their execution. When a process is created, the operating system allocates memory for its PCB and initializes it with the necessary information. The PCB is then used by the scheduler to decide which process to run next based on its state, priority, and other factors. This ensures efficient and orderly execution of processes within the operating system.

Common operations performed on a Process Control Block (PCB) include:

Process Creation

When a new process is created, the operating system allocates a PCB and initializes it with the process's initial state, PID, program counter, and other relevant information.

Process Scheduling

The scheduler uses the PCB to determine which process to run next based on its state, priority, and other scheduling criteria.

Context Switching

During a context switch, the operating system saves the current state of the CPU registers and program counter of the running process into its PCB and loads the state from the PCB of the next process to be executed.

Process Termination

When a process terminates, the operating system updates its PCB to reflect the termination state and may deallocate the PCB's memory.

Process Suspension and Resumption

The operating system can suspend a process by saving its state in the PCB and later resume it by restoring the state from the PCB.

Resource Allocation

The operating system allocates resources such as memory, CPU time, and I/O devices to the process as needed and updates the PCB to reflect these allocations.

How to Access Linux PCB from sched.h

To locate the sched.h file in the Linux source code, navigate to the following directory:

```
cd /usr/src/linux-source-5.15.0/include/linux
```

You can then open the sched.h file to examine its contents.

PCB in sched.h

The sched.h file defines the struct task_struct, which is the Linux kernel's representation of a process control block (PCB). Below is an excerpt from sched.h showing the definition of struct task_struct:

```
/* SPDX-License-Identifier: GPL-2.0 */
#ifndef _LINUX_SCHED_H
#define _LINUX_SCHED_H

/*
 * Define 'struct task_struct' and provide the main scheduler
 * APIs (schedule(), wakeup variants, etc.)
 */
```

```
struct task_struct
{
#ifdef CONFIG_THREAD_INFO_IN_TASK
    struct thread_info thread_info;
#endif
    unsigned int __state;

#ifdef CONFIG_PREEMPT_RT
    unsigned int saved_state;
#endif

    randomized_struct_fields_start

    void *stack;
    refcount_t usage;
    unsigned int flags;
    unsigned int ptrace;

#ifdef CONFIG_SMP
    int on_cpu;
    struct __call_single_node wake_entry;
#endif
#ifdef CONFIG_THREAD_INFO_IN_TASK
    unsigned int cpu;
#endif
    unsigned int wakee_flips;
    unsigned long wakee_flip_decay_ts;
    struct task_struct *last_wakee;
    int recent_used_cpu;
    int wake_cpu;
#endif
    int on_rq;

    int prio;
    int static_prio;
```

```
int normal_prio;
unsigned int rt_priority;

const struct sched_class *sched_class;
struct sched_entity se;
struct sched_rt_entity rt;
struct sched_dl_entity dl;

#ifdef CONFIG_SCHED_CORE
    struct rb_node core_node;
    unsigned long core_cookie;
    unsigned int core_occupation;
#endif

#ifdef CONFIG_CGROUP_SCHED
    struct task_group *sched_task_group;
#endif

#ifdef CONFIG_UCLAMP_TASK
    struct uclamp_se uclamp_req[UCLAMP_CNT];
    struct uclamp_se uclamp[UCLAMP_CNT];
#endif

    struct sched_statistics stats;

#ifdef CONFIG_PREEMPT_NOTIFIERS
    struct hlist_head preempt_notifiers;
#endif

#ifdef CONFIG_BLK_DEV_IO_TRACE
    unsigned int btrace_seq;
#endif

    unsigned int policy;
```

```

    int nr_cpus_allowed;
    const cpumask_t *cpus_ptr;
    cpumask_t *user_cpus_ptr;
    cpumask_t cpus_mask;
    void *migration_pending;
#ifdef CONFIG_SMP
    unsigned short migration_disabled;
#endif
    unsigned short migration_flags;

#ifdef CONFIG_PREEMPT_RCU
    int rcu_read_lock_nesting;
    union rcu_special rcu_read_unlock_special;
    struct list_head rcu_node_entry;
    struct rcu_node *rcu_blocked_node;
#endif

#ifdef CONFIG_TASKS_RCU
    unsigned long rcu_tasks_nvcsw;
    u8 rcu_tasks_holdout;
    u8 rcu_tasks_idx;
    int rcu_tasks_idle_cpu;
    struct list_head rcu_tasks_holdout_list;
#endif
};

```

This structure contains all the necessary information for managing processes, including state, priority, CPU usage, and scheduling details. It is a critical component of the Linux kernel's process management system.