

# Operating Systems


Isfahan University of Technology  
Electrical and Computer Engineering Department

Zeinab Zali

Threads Concepts and API

# A Client-Server program

Assume you have a server that is responsible for responding some clients. The clients frequently ask the server to send a requested large file. How does server manage the requests from the clients? What is the problem? What is the solution?





# Motivation


---

- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
  - Update display
  - Fetch data
  - Spell checking
  - Answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Examples of multi-thread applications: basic sorting, trees, and graph algorithms, programmers who must solve contemporary CPU-intensive problems in data mining, graphics, and artificial intelligence can leverage the power of modern multicore systems by designing solutions that run in parallel.






# Multithreaded Server Architecture






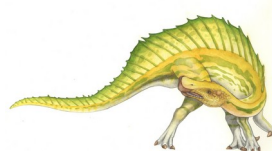
# Single and Multithreaded Processes



single-threaded process



multithreaded process





# Benefits

- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces (if the time-consuming operation is performed in a separate, asynchronous thread, the application remains responsive to the user)
- **Resource Sharing** – threads share the memory and the resources of the process to which they belong by default, so easier than shared memory or message passing between processes
- **Economy** – thread creation consumes less time and memory than process creation. Additionally, context switching is typically faster between threads than between processes
- **Scalability** – a single process can take advantage of multiprocessor architectures





# Multi-thread kernel

---


- Most operating system kernels are also typically multithreaded
- The command `ps -ef` can be used to display the kernel threads on a running Linux system
  - Examining the output of this command will show the kernel thread `kthreadd` (with `pid = 2`), which serves as the parent of all other kernel threads.






# Concurrency vs. Parallelism

- **Concurrency** supports more than one task making progress
  - Single processor / core, scheduler providing concurrency
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrent execution on single-core system:**



- **Parallelism on a multi-core system:**








# Multicore Programming (Cont.)

## ■ Types of parallelism

- **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
  - Ex: calculating an array sum or matrix multiplication
- **Task parallelism** – distributing threads across cores, each thread performing unique operation
  - Ex: calculating different statistical operation on the array of elements

## ● Hybrid





# Multicore Programming challenges

---

- **Dividing activities:** examining applications to find areas that can be divided into separate, concurrent tasks
- **Balance:** ensure that the tasks perform equal work of equal value.
- **Data splitting:** the data accessed and manipulated by the tasks must be divided to run on separate cores
- **Data dependency:** When one task depends on data from another, programmers must ensure that the execution of the tasks is synchronized to accommodate the data dependency
- **Testing and debugging:** When a program is running in parallel on multiple cores, many different execution paths are possible making debugging difficult





# User Threads and Kernel Threads

---

- **User threads** - management done by user-level threads library
- Three primary thread libraries:
  - POSIX **Pthreads**
  - Windows threads
  - Java threads
- **Kernel threads** - Supported by the Kernel
  - Examples – virtually all general purpose operating systems, including:
    - Windows
    - Solaris
    - Linux
    - Tru64 UNIX
    - Mac OS X



# تحقیق

- با بررسی در هدر فایل sched.h و جستجو، تحقیق کنید آیا برای مدیریت threadها نیز ساختاری مشابه task\_struct در سیستم عامل برای هر thread استفاده میشود؟ یا روش دیگری وجود دارد؟ تفاوتها را مشخص کنید
- مدل threadها در زبانهای C، جاوا و پایتون را مقایسه کنید.
- نحوه ساخت kernel thread و user thread با استفاده از system callهایی شبیه fork به چه صورت است؟



# Thread Libraries

---

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
  - Library entirely in user space
  - Kernel-level library supported by the OS





# Pthreads

---

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ***Specification***, not ***implementation***
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)





# Pthread.h

---

- ❑ set the thread attributes
  - `Pthread_attr_init( pthread_attr_t * attr)`
  - Examples: setting the stack size or information about the scheduling priority of the thread
- ❑ create the thread
  - `pthread_create(pthread_t *tid, pthread_attr_t *attr_t, void* thread_runner, void *thread_runner_args)`
  - `thread_runner`: function pointer
  - having a void pointer as an argument to the function `thread_runner` allows us to pass in any type of argument; having it as a return value allows the thread to return any type of result.





# Pthread.h

---

- ❑ wait for the thread to exit
  - `pthread_join(pthread_t *tid, void ** thread_runner_ret_val)`
  - The first is of type `pthread_t`, and is used to specify which thread to wait for.
  - The second argument is a pointer to the return value you expect to get back.
- ❑ terminates the calling thread
  - `pthread_exit(void * pthread_runner_ret_val)`
  - returns a value via `retval` that (if the thread is joinable) is available to another thread in the same process that calls







# Pthreads Example

---

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```





```
/* get the default attributes */
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid,&attr,runner,argv[1]);
/* wait for the thread to exit */
pthread_join(tid,NULL);

printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```





# Pthreads Code for Joining 10 Threads

---

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```





# Implicit Threading

---

- Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads
- Creation and management of threads done by compilers and run-time libraries rather than programmers
- some methods explored
  - Thread Pools
  - OpenMP
  - Fork-Join
  - Grand Central Dispatch
  - Intel Threading Building Blocks





# Client-Server Example

---

Remember our own example

What happens if the number of alive threads exceeds the maximum number of concurrent threads that the system can support?

How can we prevent this issue?





# Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
  - Usually slightly faster to service a request with an existing thread than create a new thread
  - Allows the number of threads in the application(s) to be bound to the size of the pool
  - Separating task to be performed from mechanics of creating task allows different strategies for running task
    - ▶ i.e. Tasks could be scheduled to run periodically
- Windows API supports thread pools:

```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
     * this function runs as a separate thread.  
     */  
}
```





# Sample thread pool API

---


```
void first_task() {...}  
void second_task() {...}  
void third_task() {...}  
main(){  
    // Create a thread pool.  
    pool tp(2);  
    //Add some tasks to the pool.  
    tp.schedule(&first_task);  
    tp.schedule(&second_task);  
    tp.schedule(&third_task);  
}
```





# Fork-Join Parallelism

- Multiple threads (tasks) are **forked**, and then **joined**.







# Fork-Join Parallelism

---

- General algorithm for fork-join strategy:

```
Task(problem)
  if problem is small enough
    solve the problem directly
  else
    subtask1 = fork(new Task(subset of problem))
    subtask2 = fork(new Task(subset of problem))

    result1 = join(subtask1)
    result2 = join(subtask2)

    return combined results
```







# OpenMP

- Set of compiler directives and an API for C, C++, FORTRAN provides support for parallel programming
- Identifies **parallel regions** – blocks of code that can run in parallel

```
#pragma omp parallel
```

Create as many threads as there are cores

```
#pragma omp parallel for  
for(i=0;i<N;i++) {  
    c[i] = a[i] + b[i];  
}
```

Run for loop in parallel





# Threading Issues

---

- Semantics of **fork()** and **exec()** system calls
- Signal handling
  - Synchronous and asynchronous
- Thread cancellation of target thread
  - Asynchronous or deferred





# Semantics of `fork()` and `exec()`

---

- Does `fork()` duplicate only the calling thread or all threads?
  - Some UNIX systems have two versions of `fork()`:
    - one that duplicates all threads
    - one that duplicates only the thread that invoked the `fork()` system call.
- `exec()` usually works as normal – replace the running process including all threads
  - So when `exec()` is called immediately after forking, forking only the calling thread is sufficient





# Signal Handling

- **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- A **signal handler** is used to process signals
  1. Signal is generated by particular event
  2. Signal is delivered to a process
  3. Signal is handled by one of two signal handlers:
    1. default
    2. user-defined
- Every signal has **default handler** that kernel runs when handling signal
  - **User-defined signal handler** can override default
  - For single-threaded, signal delivered to process





# Signal Handling (Cont.)

---

- Where should a signal be delivered for multi-threaded?
  - Deliver the signal to the thread to which the signal applies
  - Deliver the signal to every thread in the process
  - Deliver the signal to certain threads in the process
  - Assign a specific thread to receive all signals for the processes

**kill(pid\_t, signal)**

**pthread\_kill(thread\_t, signal)**





# Thread Cancellation

- Terminating a thread before it has finished
  - Suppose multiple threads searching for a record in a database and one of them find it
  - a web page loads using several threads—each image is loaded in a separate thread. When a user presses the stop button on the browser, all threads loading the page are canceled.
- Thread to be canceled is **target thread**
- Two general approaches:
  - **Asynchronous cancellation** terminates the target thread immediately
  - **Deferred cancellation** allows the target thread to periodically check if it should be canceled







# Thread Cancellation (Cont.)

- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation disabled, cancellation remains pending until thread enables it
- Default type is deferred
  - Cancellation only occurs when thread reaches **cancellation point**
    - ▶ Ex: `pthread_testcancel()`
    - ▶ Ex: `read` function
- On Linux systems, thread cancellation is handled through signals





# Thread-Local Storage

---

- **Thread-local storage (TLS)** allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- Different from local variables
  - Local variables visible only during single function invocation
  - TLS visible across function invocations
- Similar to `static` data
  - TLS is unique to each thread





# Linux Threads

- Linux refers to them as **tasks** rather than **threads**
- Thread creation is done through `clone()` system call
- `clone()` allows a child task to share the address space of the parent task (process)
  - Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- `struct task_struct` points to process data structures (shared or unique)

getconf GNU\_LIBPTHREAD\_VERSION





# End of Chapter 4

